

UN Data Hackathon

extracting AIS data from UNGP

Team : GEMy

group email :

gemy@dosm.gov.my (<mailto:gemy@dosm.gov.my>).

individual email:

rajkumar@dosm.gov.my (<mailto:rajkumar@dosm.gov.my>) ; shukor.talib@dosm.gov.my (<mailto:shukor.talib@dosm.gov.my>) ; nurhuda@dosm.gov.my (<mailto:nurhuda@dosm.gov.my>) ; tgnoradilah@dosm.gov.my (<mailto:tgnoradilah@dosm.gov.my>) ; najmi.ariffin@dosm.gov.my (<mailto:najmi.ariffin@dosm.gov.my>).

In [1]:

```
#allow multiple outputs in one jupyter cell
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

import pandas as pd
from datetime import datetime
# to apply aggregation functions on spark df
import pyspark.sql.functions as F
```

In [2]:

```
# this cell contains the code to access GitLab repo
# need it to install ais package from GitLab repo
import sys
import subprocess

GITLAB_USER = "read_aistt" # read only access
GITLAB_TOKEN = "MMQ6ky1rnLsuKxjyZuvB"

# clone the repo and install the ais packag
git_package = f"git+https://{GITLAB_USER}:{GITLAB_TOKEN}@code.officialstatistics.org/trade-

std_out = subprocess.run([sys.executable, "-m", "pip", "install", git_package], capture_out
print(std_out)
```

```
Collecting git+https://read_aistt:****@code.officialstatistics.org/trade-tas
k-team-phase-1/ais.git
  Cloning https://read_aistt:****@code.officialstatistics.org/trade-task-tea
m-phase-1/ais.git (https://read_aistt:****@code.officialstatistics.org/trade
-task-team-phase-1/ais.git) to /tmp/pip-req-build-3p5s810y
Building wheels for collected packages: ais
  Building wheel for ais (setup.py): started
  Building wheel for ais (setup.py): finished with status 'done'
  Created wheel for ais: filename=ais-2.7.6-py3-none-any.whl size=9267 sha25
6=050e877075184ea051f78f61665e66484c2bf6c3165c83f0eaa71e002df07a0c
  Stored in directory: /tmp/pip-ephem-wheel-cache-fjpleh21/wheels/49/e0/a2/2
5d96a62cf626776ab2fd57fcbd822c2b8118049a84b16953d
Successfully built ais
Installing collected packages: ais
Successfully installed ais-2.7.6
```

In [3]:

```
# import get_ais() from ais package
from ais import functions as af
```

In [4]:

```
# details about the function e.g.
# input parameters,
# output: spark df
# usage of this function in examples below
af.get_ais?
```

Signature:

```
af.get_ais(
    spark: pyspark.sql.session.SparkSession,
    start_date: datetime.datetime,
    end_date: datetime.datetime = None,
    h3_list: Union[List[int], NoneType] = None,
    polygon_hex_df: Union[pyspark.sql.dataframe.DataFrame, NoneType] = None,
    mmsi_list: Union[List[int], NoneType] = None,
    message_type: Union[List[int], NoneType] = [1, 2, 3, 4, 18, 19, 27],
    columns: Union[List[str], NoneType] = ['*'],
    polygon: Union[Dict, NoneType] = None,
    polygon_hex_resolution: Union[int, NoneType] = 8,
) -> pyspark.sql.dataframe.DataFrame
```

Docstring:

A wrapper function to apply filters on the AIS data.
Note that default parameters for message type are
position message types

Parameters

spark: SparkSession

start_date: datetime
the start date filter to apply

end_date: datetime
the end date filter to apply. To filter a single date, use end_date equal to start_date

h3_list: list of int, default None
h3 indices must be in int format and must have the same resolution.
if None then it is not applied.

polygon_hex_df: dataframe, from polygon_to_hex_df function
Dataframe with the following columns (minimum columns to contain) :
- hex_id: the h3 hex ids (64-bit ints)
- polygon_name: the name of the polygon
- hex_resolution: the resolution of the hex (should be the same for all)
The hex_ids should be contained in only one polygon_name, otherwise resulting dataframe
will contain duplicate entries.

mmsi_list: list of int, default None
the list of mmsi filter to apply. if None, then it is not applied

message_type: list of int, default [1,2,3,4,18,19,27] <- position messages
the list of message types to retain. if not supplied then the default message type filter is applied
use ["*"] to get all message types

columns: list of str, default ["*"]
the list of columns to retain. if not supplied, all columns are returned

`polygon: Optional[Dict] = None`

GeoJson representation of polygon. If supplied, then the hex approximation of the polygon is calculated using `poly_container` (`polygon`, `hex_resolution`, `overflow=True`). The AIS data will be filtered according to hexes first and then according to polygon using Sedona functions:

```
select *, ST_Point(longitude,latitude) as point, ST_GeomFromGeoJSON('{polygon}') as
polygon from temp where ST_Within(point, polygon)
```

`polygon_hex_resolution: int = 8`

The resolution of the hexagons to fill the input polygon with. Default is 8, a hex with an avg area of 0.737 sq km.

A polygon with an area of 100 sq. km will contain ~136 resolution 8 hexes. The same 100 sq. km polygon

can be approximated by ~949 hexes using resolution 9. Note that the higher the resolution, the higher

the polygon area covered by the hexes. However, a small increase in resolution dramatically increases

the number of hexes. See <https://h3geo.org/docs/core-library/restable/> (<https://h3geo.org/docs/core-library/restable/>) for a table of hex resolutions.

Returns

Spark dataframe with the filters applied.

Notes

If multiple filters are provided, the most restrictive filters are applied.

For example, both `polygon` and `h3_list`

are provided where `h3_list` is a list of hexes fully contained within the polygon. The filtered AIS data will only contain

those within the hexes. Data within the polygon but outside the hexes will not be included.

File: /opt/conda/lib/python3.8/site-packages/ais/_aisfilter.py

Type: function

In [5]:

```
start_date = datetime.fromisoformat("2022-01-01")
```

```
df = af.get_ais(spark, start_date)
df.show(n=1, vertical=True, truncate=False)
```

```
-RECORD 0-----
-----
message_type      | 1
mmsi              | 205654000
imo              | 9691279
vessel_name       | DN97
callsign          | ORRK
vessel_type       | Port Tender
vessel_type_code  | 53
vessel_type_cargo | null
vessel_class      | A
length           | 17.0
width            | 6.0
flag_country      | Belgium
flag_code         | 205
destination       | ZEEBRUGGE
eta              | 10081400
draught          | 1.2
longitude         | 3.20316
latitude          | 51.32248833
sog              | 0.0
cog              | 0.0
rot              | 0.0
heading          | 0.0
nav_status        | Not Defined
nav_status_code   | 15
source            | T-AIS
dt_pos_utc        | 2022-01-01 21:17:43
dt_static_utc     | 2022-01-01 21:14:24
dt_insert_utc     | 2022-01-01 21:17:48
vessel_type_main  | null
vessel_type_sub   | null
eeid              | 4897682788452534256
source_filename   | s3a://ungp-ais-data-historical-backup/exact-earth-data/
nonprod/year=2022/month=01/day=01/20220101211833.csv.gz
H3index_0         | 8019ffffffffffff
H3_int_index_0    | 576918149140578303
H3_int_index_1    | 581412952674926591
H3_int_index_2    | 585913253767413759
H3_int_index_3    | 590416715955830783
H3_int_index_4    | 594920306993266687
H3_int_index_5    | 599423900178186239
H3_int_index_6    | 603927499402903551
H3_int_index_7    | 608431098929610751
H3_int_index_8    | 612934698542301183
H3_int_index_9    | 617438298168098815
H3_int_index_10   | 621941897795403775
H3_int_index_11   | 626445497422761983
H3_int_index_12   | 630949097050129919
H3_int_index_13   | 635452696677454463
H3_int_index_14   | 639956296304824911
H3_int_index_15   | 644459895932195403
only showing top 1 row
```

In [6]:

```
#malaysia
columns = ["mmsi", "latitude", "longitude", "dt_insert_utc", "eeid", "flag_country", "flag_
            "eta", "vessel_type_main", "imo", "vessel_type", "vessel_type_code", "vessel_type_
            "destination", "vessel_type_sub" ]

start_date = datetime.fromisoformat("2022-01-01")
end_date = datetime.fromisoformat("2022-01-31")
df = af.get_ais(spark,
                start_date,
                end_date = end_date,
                columns=columns)

df.count()
```

Out[6]:

708862027

In [7]:

```
# filter mmsi

mmsi_list = [533131111,
             533131162,
             311053500,
             357900000,
             370220000,
             209444000,
             235479000,
             249675000,
             255805778,
             309046000,
             311053500,
             312067000,
             370220000,
             374898000,
             477311600,
             565711000,
             413212060 ]

columns = ["mmsi", "latitude", "longitude", "dt_insert_utc", "eeid", "flag_country", "flag_
            "eta", "vessel_type_main", "imo", "vessel_type", "vessel_type_code", "vessel_type_
            "destination", "vessel_type_sub" ]

start_date = datetime.fromisoformat("2022-01-01")
end_date = datetime.fromisoformat("2022-01-31")

df = af.get_ais(spark,
                start_date,
                end_date = end_date,
                mmsi_list = mmsi_list,
                columns = columns)

df.count()
```

Out[7]:

93140

In [8]:

```
df.show(n=1, vertical=True, truncate=False)
```

```
-RECORD 0-----  
--  
mmsi           | 235479000  
latitude       | 5.86358667  
longitude      | 96.49673167  
dt_insert_utc  | 2022-01-01 22:51:48  
eeid           | 5115565929652625368  
flag_country   | UK  
flag_code      | 235  
eta            | 1080600  
vessel_type_main | Container Ship  
imo            | 9241322  
vessel_type    | Cargo  
vessel_type_code | 74  
vessel_type_cargo | Carrying DG,HS or MP,IMO hazard or Pollutant Category 0  
S  
destination    | MYPKG=>AEJBA  
vessel_type_sub | null  
only showing top 1 row
```



In [9]:

```
# first this function and then pass on its output with get_ais()
af.polygon_to_hex_df?
```

Signature:

```
af.polygon_to_hex_df(
    polygons: List[Tuple[str, Dict]],
    hex_resolution: int = 8,
    overfill=False,
) -> pandas.core.frame.DataFrame
```

Docstring:

A wrapper for h3.polyfill that returns integer hex ids for multiple polygons.

Parameters

polygons: list of tuples

the first element in this tuple is expected to be a (name) string identifier for the polygon and the second element is the polygon itself (see example above)

hex_resolution: int, default 8

the resolution of the hexagons to fill the input polygon with. Default is 8, a hex with an avg area of 0.737 sq km.

A polygon with an area of 100 sq. km will contain ~136 resolution 8 hexes. The same 100 sq. km polygon

can be approximated by ~949 hexes using resolution 9. Note that the higher the resolution, the higher

the polygon area covered by the hexes. However, a small increase in resolution dramatically increases

the number of hexes. See <https://h3geo.org/docs/core-library/restable/> (<https://h3geo.org/docs/core-library/restable/>) for a table of hex resolutions.

Returns

Dataframe with the following columns:

- hex_id: the h3 hex ids (64-bit ints)
- polygon_name: the name of the polygon
- hex_resolution: the resolution of the hex

```
>>>hull_bbox = {
    "type": "Polygon",
    "coordinates": [
        [
            [-0.3169, 53.7344],
            [-0.2537, 53.7344],
            [-0.2537, 53.75],
            [-0.3169, 53.75],
            [-0.3169, 53.7344]
        ]
    ]
}
>>>london_bbox = {
    "type": "Polygon",
    "coordinates": [
        [
            [-0.1203, 51.4415],
            [0.5869, 51.4415],
            [0.5869, 51.5262],
```



```

        [-0.1203, 51.5262],
        [-0.1203, 51.4415]
    ]
}
>>>query_polys = [("HullPortArea", hull_bbox),("LondonPortArea", london_bbox)]
>>>polygon_to_hex_df(query_polys, 10)

```

	hex_id	polygon_name	hex_resolution
0	621940969126789119	HullPortArea	10
1	621940969212772351	HullPortArea	10
2	621940969214869503	HullPortArea	10
3	621940974334017535	HullPortArea	10
4	621940969137274879	HullPortArea	10

34427	621941942979756031	LondonPortArea	10
34428	621941942943285247	LondonPortArea	10
34429	621941942854713343	LondonPortArea	10
34430	621941942818242559	LondonPortArea	10
34431	621941940535295999	LondonPortArea	10

```

File:      /opt/conda/lib/python3.8/site-packages/ais/_poly.py
Type:      function

```

In [10]:

```

# 2nd parameter for polygon_to_hex_df()
# https://boundingbox.klokantech.com/
# polygon coordinates in geojson format
malaysia_polygon = {
    "type": "Polygon",
    "coordinates": [
        [
            [101.1134319752,2.8025194725],
            [101.4495449513,2.8025194725],
            [101.4495449513,3.1131551166],
            [101.1134319752,3.1131551166],
            [101.1134319752,2.8025194725]
        ]
    ]
}

```



In [11]:

```
# first parameter for polygon_to_hex_df() is the name/label for the polygon
polygon_hex_df_malaysia = af.polygon_to_hex_df([("Malaysia_Port_Polygon", malaysia_polygon)])
```

In [12]:

```
start_date = datetime.fromisoformat("2022-01-01")
end_date = datetime.fromisoformat("2022-01-31")
columns = ["mmsi", "latitude", "longitude", "dt_insert_utc", "eeid", "flag_country", "flag_
            "eta", "vessel_type_main", "imo", "vessel_type", "vessel_type_code", "vessel_type_
            "destination", "vessel_type_sub" ]

# pass polygon_hex_df to get_ais()
df = af.get_ais(spark,
                start_date,
                end_date = end_date,
                columns = columns,
                polygon_hex_df = polygon_hex_df_malaysia
                )

df.count()
```

Out[12]:

797879

In [13]:

```
mmsi_list = [533131111,
             533131162,
             311053500,
             357900000,
             370220000,
             209444000,
             235479000,
             249675000,
             255805778,
             309046000,
             311053500,
             312067000,
             370220000,
             374898000,
             477311600,
             565711000,
             413212060 ]

start_date = datetime.fromisoformat("2022-01-01")
end_date = datetime.fromisoformat("2022-01-31")
columns = ["mmsi", "latitude", "longitude", "dt_insert_utc", "eeid",
           "flag_country", "flag_code",
           "eta", "vessel_type_main", "imo", "vessel_type", "vessel_type_code", "vessel_type_
           "destination", "vessel_type_sub" ]

# pass polygon_hex_df to get_ais()
df = af.get_ais(spark,
                start_date,
                end_date = end_date,
                columns = columns,
                mmsi_list = mmsi_list,
                polygon_hex_df = polygon_hex_df_malaysia

                )

df.count()
```

Out[13]:

20701

In [14]:

```
# ais messages captured in the Malaysia port region
df.show(n=1)
```

```
+-----+-----+-----+-----+-----+-----+-----+
|hex_resolution| longitude| mmsi|flag_country| eta|destination|vess
el_type_code|vessel_type_cargo|vessel_type| eeid|flag_code|
imo| H3_int_index_8| latitude|vessel_type_main| dt_insert_utc|vesse
l_type_sub| polygon_name|
+-----+-----+-----+-----+-----+-----+-----+
|
8|101.30643167|209444000| Cyprus|1011600| MYPKG|
70| null| Cargo|4705970841525673120| 209|9507714|61426
6716666462207|2.80780333| null|2022-01-01 17:42:21| nul
l|Malaysia_Port_Pol...|
+-----+-----+-----+-----+-----+-----+-----+
only showing top 1 row
```

In [15]:

```
pd_df = df.toPandas()

type(pd_df)
pd_df.shape
```

Out[15]:

pandas.core.frame.DataFrame

Out[15]:

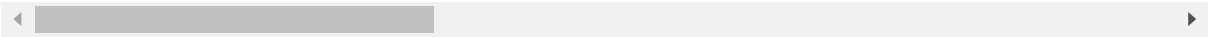
(20701, 18)

In [16]:

```
pd_df.head()
```

Out[16]:

	hex_resolution	longitude	mmsi	flag_country	eta	destination	vessel_type_code
0	8	101.306432	209444000	Cyprus	1011600	MYPKG	70
1	8	101.309188	209444000	Cyprus	1011600	MYPKG	70
2	8	101.306525	209444000	Cyprus	1011600	MYPKG	70
3	8	101.306937	209444000	Cyprus	1011600	MYPKG	70
4	8	101.306852	209444000	Cyprus	1011600	MYPKG	70



In [17]:

```
!pip install s3fs
import s3fs

# create a handle for s3fs
fs = s3fs.S3FileSystem(anon=False)
```

WARNING: The directory '/home/sparkuser/.cache/pip' or its parent directory is not owned or is not writable by the current user. The cache has been disabled. Check the permissions and owner of that directory. If executing pip with sudo, you should use sudo's -H flag.

Collecting s3fs

Downloading s3fs-2022.11.0-py3-none-any.whl (27 kB)

Collecting fsspec==2022.11.0

Downloading fsspec-2022.11.0-py3-none-any.whl (139 kB)

|██| 139 kB 52.3 MB/s eta 0:00:01

Collecting aiohttp!=4.0.0a0,!4.0.0a1

Downloading aiohttp-3.8.3-cp38-cp38-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1.0 MB)

|██| 1.0 MB 128.7 MB/s eta 0:00:01

Collecting aiobotocore~=2.4.0

Downloading aiobotocore-2.4.0-py3-none-any.whl (65 kB)

|██| 65 kB 98.6 MB/s eta 0:00:01

Collecting boto3<1.27.60,>=1.27.59

Downloading boto3-1.27.59-py3-none-any.whl (9.1 MB)

|██| 9.1 MB 99.7 MB/s eta 0:00:01

Collecting aioitertools>=0.5.1

Downloading aioitertools-0.11.0-py3-none-any.whl (23 kB)

Collecting wrapt>=1.10.10

Downloading wrapt-1.14.1-cp38-cp38-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_17_x86_64.manylinux2014_x86_64.whl (81 kB)

|██| 81 kB 95.3 MB/s eta 0:00:01

Collecting yarll<2.0,>=1.0

Downloading yarll-1.8.1-cp38-cp38-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (262 kB)

|██| 262 kB 129.8 MB/s eta 0:00:01

Collecting async-timeout<5.0,>=4.0.0a3

Downloading async_timeout-4.0.2-py3-none-any.whl (5.8 kB)

Collecting multidict<7.0,>=4.5

Downloading multidict-6.0.2-cp38-cp38-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (121 kB)

|██| 121 kB 125.9 MB/s eta 0:00:01

Collecting frozenlist>=1.1.1

Downloading frozenlist-1.3.3-cp38-cp38-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_17_x86_64.manylinux2014_x86_64.whl (161 kB)

|██| 161 kB 118.3 MB/s eta 0:00:01

Collecting charset-normalizer<3.0,>=2.0

Downloading charset_normalizer-2.1.1-py3-none-any.whl (39 kB)

Requirement already satisfied: attrs>=17.3.0 in /opt/conda/lib/python3.8/site-packages (from aiohttp!=4.0.0a0,!4.0.0a1->s3fs) (21.2.0)

Collecting aiosignal>=1.1.2

Downloading aiosignal-1.3.1-py3-none-any.whl (7.6 kB)

Collecting typing_extensions>=4.0

Downloading typing_extensions-4.4.0-py3-none-any.whl (26 kB)

Collecting jmespath<2.0.0,>=0.7.1

Downloading jmespath-1.0.1-py3-none-any.whl (20 kB)

Requirement already satisfied: urllib3<1.27,>=1.25.4 in /opt/conda/lib/python3.8/site-packages (from boto3<1.27.60,>=1.27.59->aiobotocore~=2.4.0->s3fs) (1.26.4)

Requirement already satisfied: python-dateutil<3.0.0,>=2.1 in /opt/conda/li

